

# Examples and Exercises from Think Stats, 2nd Edition

<http://thinkstats2.com>

Copyright 2016 Allen B. Downey

MIT License: <https://opensource.org/licenses/MIT>

## Exercises

Using data from the NSFG, make a scatter plot of birth weight versus mother's age. Plot percentiles of birth weight versus mother's age. Compute Pearson's and Spearman's correlations. How would you characterize the relationship between these variables?

```
In [28]: import first

live, firsts, others = first.MakeFrames()
live = live.dropna(subset=['agepreg', 'totalwgt_lb'])
```

```
In [29]: ages = live.agepreg
weights = live.totalwgt_lb
print('Corr', Corr(ages, weights))
print('SpearmanCorr', SpearmanCorr(ages, weights))
```

```
Corr 0.0688339703541091
SpearmanCorr 0.09461004109658226
```

```
In [30]: def BinnedPercentiles(df):
    """Bin the data by age and plot percentiles of weight for each bin.

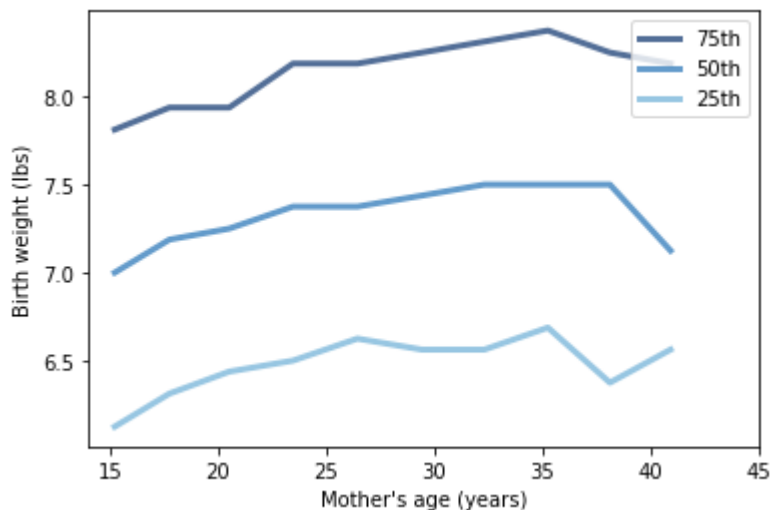
    df: DataFrame
    """
    bins = np.arange(10, 48, 3)
    indices = np.digitize(df.agepreg, bins)
    groups = df.groupby(indices)

    ages = [group.agepreg.mean() for i, group in groups][1:-1]
    cdfs = [thinkstats2.Cdf(group.totalwgt_lb) for i, group in groups][1:-1]

    thinkplot.PrePlot(3)
    for percent in [75, 50, 25]:
        weights = [cdf.Percentile(percent) for cdf in cdfs]
        label = '%dth' % percent
        thinkplot.Plot(ages, weights, label=label)

    thinkplot.Config(xlabel="Mother's age (years)",
                     ylabel='Birth weight (lbs)',
                     xlim=[14, 45], legend=True)


    BinnedPercentiles(live)
```



```
In [31]: def ScatterPlot(ages, weights, alpha=1.0, s=20):
        """Make a scatter plot and save it.

        ages: sequence of float
        weights: sequence of float
        alpha: float
        """
        thinkplot.Scatter(ages, weights, alpha=alpha)
        thinkplot.Config(xlabel='Age (years)',
                          ylabel='Birth weight (lbs)',
                          xlim=[10, 45],
                          ylim=[0, 15],
                          legend=False)

        ScatterPlot(ages, weights, alpha=0.05, s=10)
```



```
In [32]: #Answers:

# 1) It is difficult to see, but the scatterplot appears to show a weak relationship be

# 2) This is supported by the correlation. Spearman's correlation is ~ 0.09. Pearson's
# variance means either a non-linear relationship or some influence of outliers

# 3) The relationship appears to have a non-linear relationship if you plot % of weight
# If the mother is between ages of 15 and 25, the birth weight increases more quickly
# effect is weaker
```



# Examples and Exercises from Think Stats, 2nd Edition

<http://thinkstats2.com> (<http://thinkstats2.com>)

Copyright 2016 Allen B. Downey

MIT License: <https://opensource.org/licenses/MIT> (<https://opensource.org/licenses/MIT>)

## Exercises

**Exercise:** In this chapter we used  $\bar{x}$  and median to estimate  $\mu$ , and found that  $\bar{x}$  yields lower MSE. Also, we used  $S^2$  and  $S_{n-1}^2$  to estimate  $\sigma$ , and found that  $S^2$  is biased and  $S_{n-1}^2$  unbiased. Run similar experiments to see if  $\bar{x}$  and median are biased estimates of  $\mu$ . Also check whether  $S^2$  or  $S_{n-1}^2$  yields a lower MSE.

```
In [12]: ▶ def Estimate4(n=7, iters=100000):
    """Mean error for xbar and median as estimators of population mean.

    n: sample size
    iters: number of iterations
    """
    mu = 0
    sigma = 1

    means = []
    medians = []
    for _ in range(iters):
        xs = [random.gauss(mu, sigma) for i in range(n)]
        xbar = np.mean(xs)
        median = np.median(xs)
        means.append(xbar)
        medians.append(median)

    print('Experiment 1')
    print('mean error xbar', MeanError(means, mu))
    print('mean error median', MeanError(medians, mu))

Estimate4()
```

```
Experiment 1
mean error xbar -0.00020005658318784632
mean error median -0.00040362310046356225
```

```
In [13]: ▶ def Estimate5(n=7, iters=100000):
    """RMSE for biased and unbiased estimators of population variance.

    n: sample size
    iters: number of iterations
    """

    mu = 0
    sigma = 1

    estimates1 = []
    estimates2 = []
    for _ in range(iters):
        xs = [random.gauss(mu, sigma) for i in range(n)]
        biased = np.var(xs)
        unbiased = np.var(xs, ddof=1)
        estimates1.append(biased)
        estimates2.append(unbiased)

    print('Experiment 2')
    print('RMSE biased', RMSE(estimates1, sigma**2))
    print('RMSE unbiased', RMSE(estimates2, sigma**2))

Estimate5()
```

```
Experiment 2
RMSE biased 0.5133404751902225
RMSE unbiased 0.5758719020066496
```

```
In [14]: ▶ # Answers:

# 1) the xbar and the median yield lower mean error as m increases, so neither
# one is obviously biased based on this experience

# 2) The biased estimator of variance yields lower RMSE than the unbiased
# estimator approximately by 10% and the difference holds up as m increases
```

**Exercise:** In games like hockey and soccer, the time between goals is roughly exponential. So you could estimate a team's goal-scoring rate by observing the number of goals they score in a game. This estimation process is a little different from sampling the time between goals, so let's see how it works.

Write a function that takes a goal-scoring rate,  $\lambda$ , in goals per game, and simulates a game by generating the time between goals until the total time exceeds 1 game, then returns the number of goals scored.

Write another function that simulates many games, stores the estimates of  $\lambda$ , then computes their mean error and RMSE.

Is this way of making an estimate biased?

```
In [15]: ▶ def SimulateGame(lam):  
    """Simulates a game and returns the estimated goal-scoring rate.  
  
    lam: actual goal scoring rate in goals per game  
    """  
    goals = 0  
    t = 0  
    while True:  
        time_between_goals = random.expovariate(lam)  
        t += time_between_goals  
        if t > 1:  
            break  
        goals += 1  
  
    # approx goal-scoring rate = actual number of goals scored  
    L = goals  
    return L
```

```
In [16]: ▶ def Estimate6(lam=2, m=1000000):


    estimates = []
    for i in range(m):
        L = SimulateGame(lam)
        estimates.append(L)

    print('Experiment 4')
    print('rmse L', RMSE(estimates, lam))
    print('mean error L', MeanError(estimates, lam))

    pmf = thinkstats2.Pmf(estimates)
    thinkplot.Hist(pmf)
    thinkplot.Config(xlabel='Goals scored', ylabel='PMF')

Estimate6()
```

Experiment 4  
rmse L 1.4150922231430714  
mean error L 0.00059



```
In [17]: ▶ # Answers

# 1) When estimating lambda in this way, RMSE is 1.4

# 2) This estimator appears to be unbiased as the mean error decreases with m
```